

# QEMU in the Fast Lane: Accelerating KubeVirt Networking with eBPF

Daniel Borkmann, Anton Protopopov, Jussi Mäki  
KubeCon North America 2025



“In late 2023, a long-trusted virtualization staple became the biggest open question on the enterprise IT roadmap.” - MIT Technology Review

**COMPUTING**

# Turning migration into modernization

The VMware shake-up has led to an IT inflection point. Leaders are now weighing whether to renew, migrate, or redesign entirely for the cloud era.

By MIT Technology Review Insights

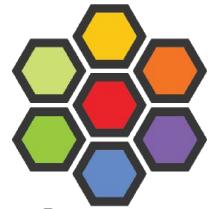
October 2, 2025

Are there viable OSS solutions which help with  
a strategic reset to modernize  
VM infrastructure into Kubernetes?

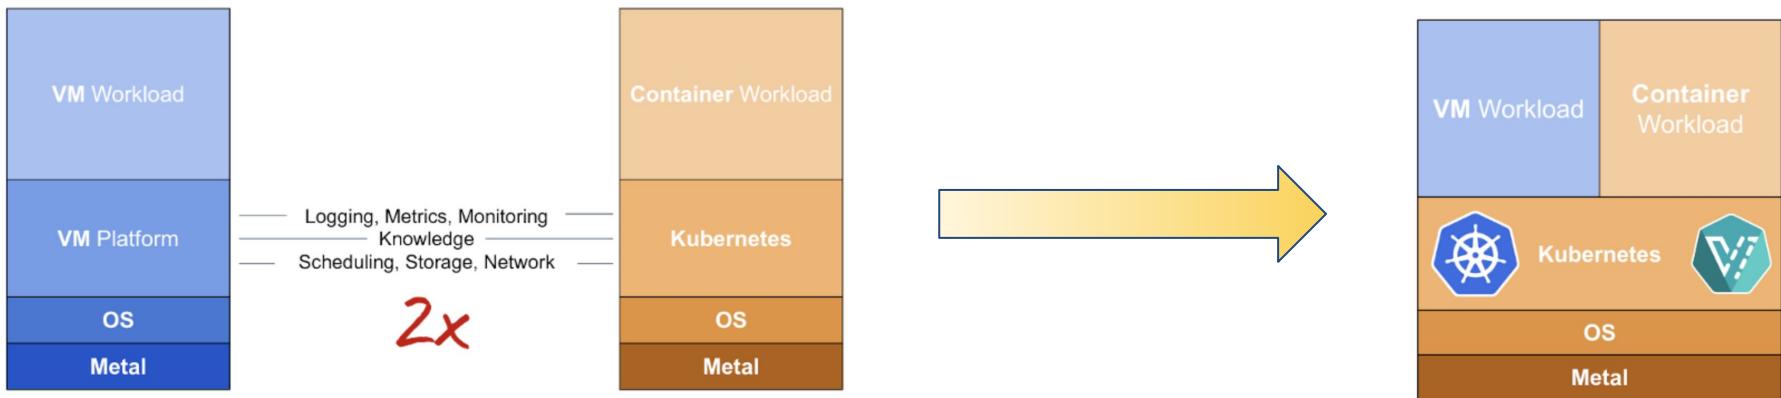
Yes! The answer is:

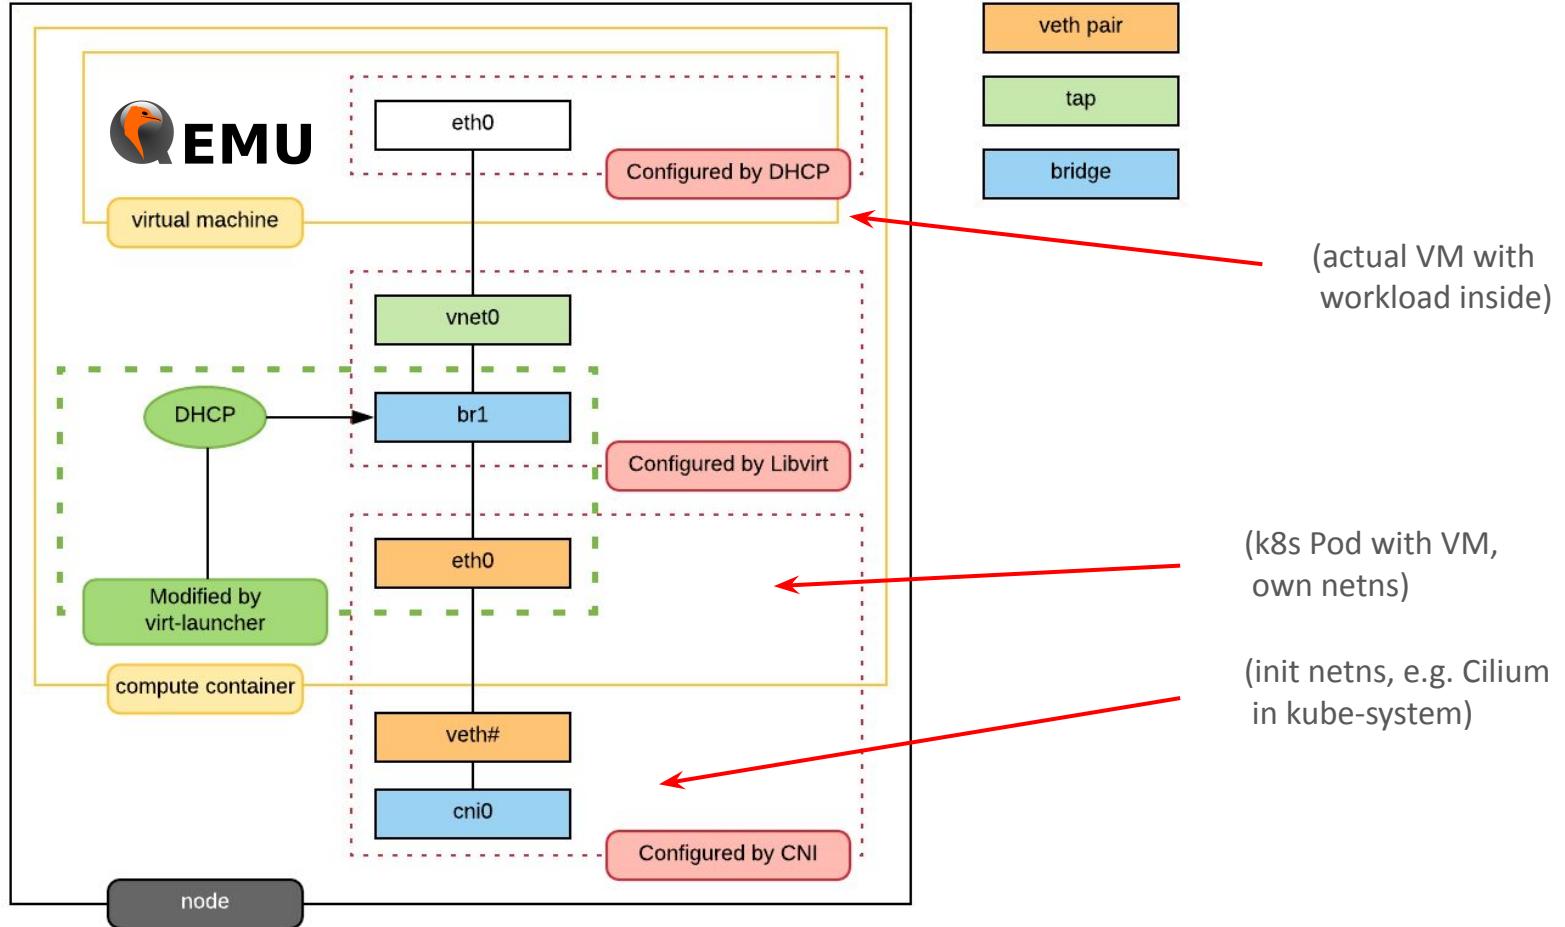


KubeVirt +

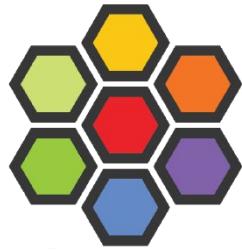


# Kubernetes & KubeVirt enables cost savings through infrastructure convergence:





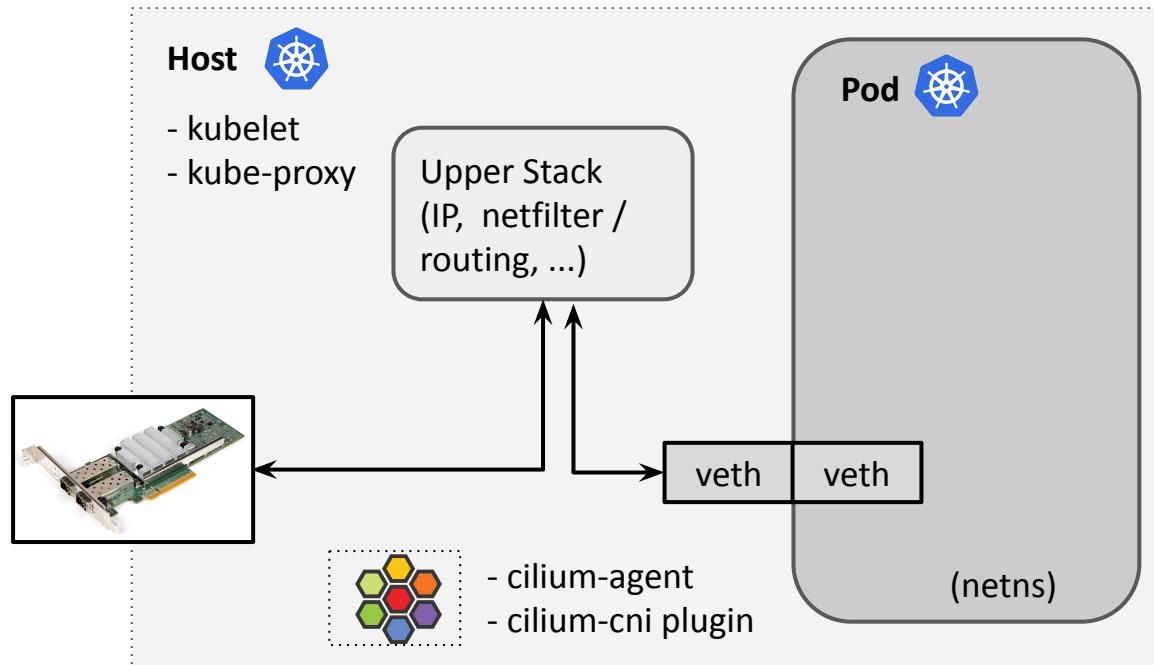
Moonshot idea for *this* talk:  
Can we reimagine how we do  
VM networking in Kubernetes?



Lets first go back to fundamentals on how we removed the Pod (netns) overhead in Cilium ...



# Standard Pod Datapath:





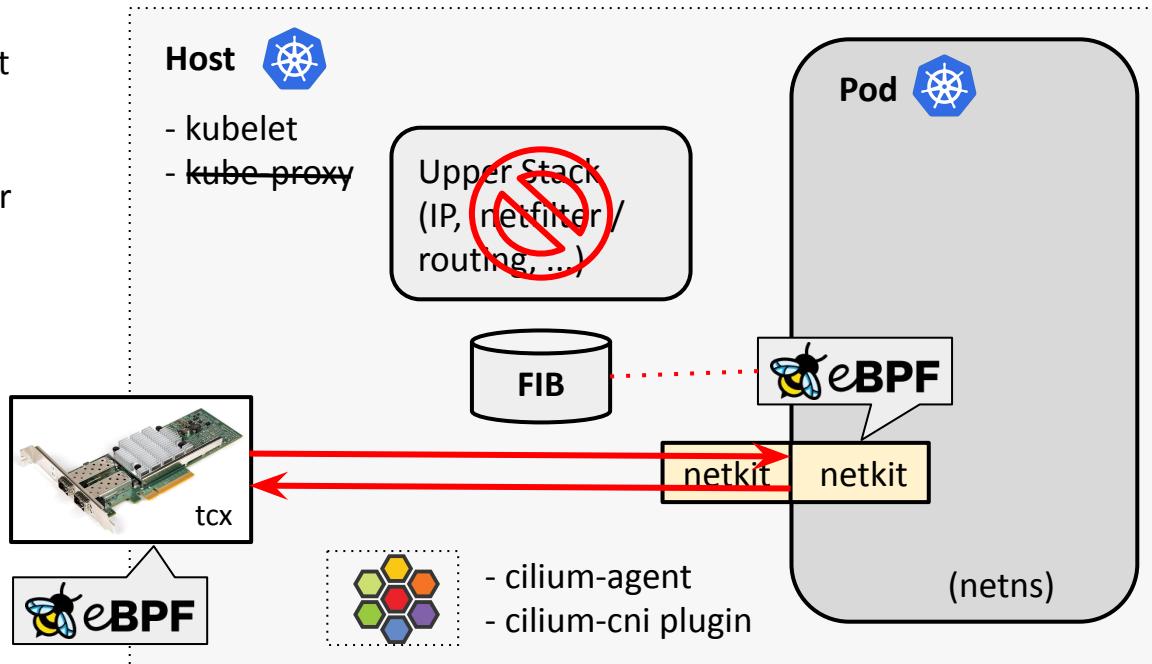
# Cilium's Pod Datapath: Making Pods as fast as Host

## Building Blocks:

- BPF kube-proxy replacement
- BPF Host Routing
- BIG TCP for GRO/GSO/TSO
- tcx-based BPF datapath layer
- netkit devices for Pods

## Recommended:

- Cilium v1.18+
- Linux kernel v6.12+





# Deep Dive: veth-replacement for Pods

**netkit programmable virtual devices for BPF:**

- Realizing multiple core ideas tailored for Cilium:
- Better application scheduler accounting by removing heavy reliance on ksoftirqd-deferral under pressure



# Deep Dive: veth-replacement for Pods

**netkit programmable virtual devices for BPF:**

- Realizing multiple core ideas tailored for Cilium:
  - Better application scheduler accounting by removing heavy reliance on ksoftirqd-deferral under pressure
  - No driver internal queueing for traffic in/out of the node, thus complete removal of netns overhead



# Deep Dive: veth-replacement for Pods

## netkit programmable virtual devices for BPF:

- Realizing multiple core ideas tailored for Cilium:
  - Better application scheduler accounting by removing heavy reliance on ksoftirqd-deferral under pressure
  - No driver internal queueing for traffic in/out of the node, thus complete removal of netns overhead
  - veth is L2, provide users choice between L3 (noarp) or L2 device option



# Deep Dive: veth-replacement for Pods

## netkit programmable virtual devices for BPF:

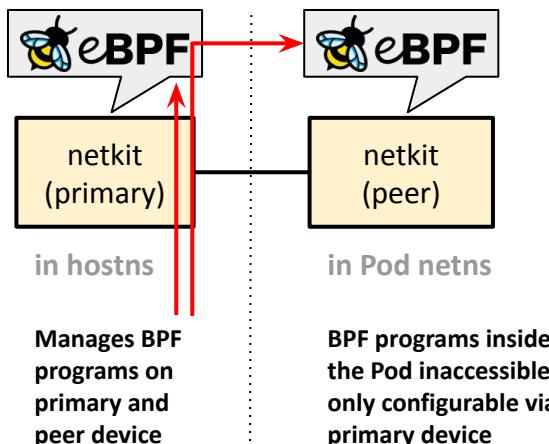
- Realizing multiple core ideas tailored for Cilium:
  - Better application scheduler accounting by removing heavy reliance on ksoftirqd-deferral under pressure
  - No driver internal queueing for traffic in/out of the node, thus complete removal of netns overhead
  - veth is L2, provide users choice between L3 (noarp) or L2 device option
  - Device pair is partitioned into primary (hostns) and peer (Pod-netns) device for BPF management



# Deep Dive: veth-replacement for Pods

## netkit programmable virtual devices for BPF:

- Realizing multiple core ideas tailored for Cilium:
- Better application scheduler accounting by removing heavy reliance on ksoftirqd-deferral under pressure
- No driver internal queueing for traffic in/out of the node, thus complete removal of netns overhead
- veth is L2, provide users choice between L3 (noarp) or L2 device option
- Device pair is partitioned into primary (hostns) and peer (Pod-netns) device for BPF management

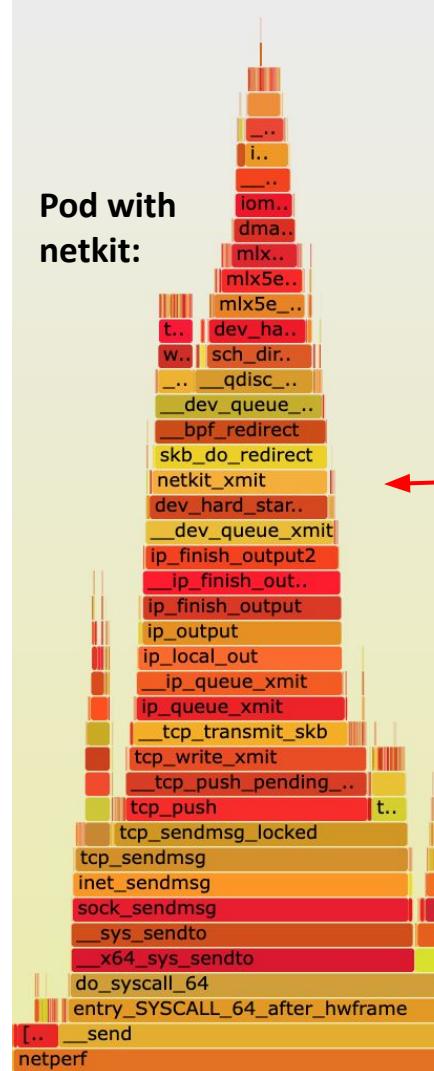
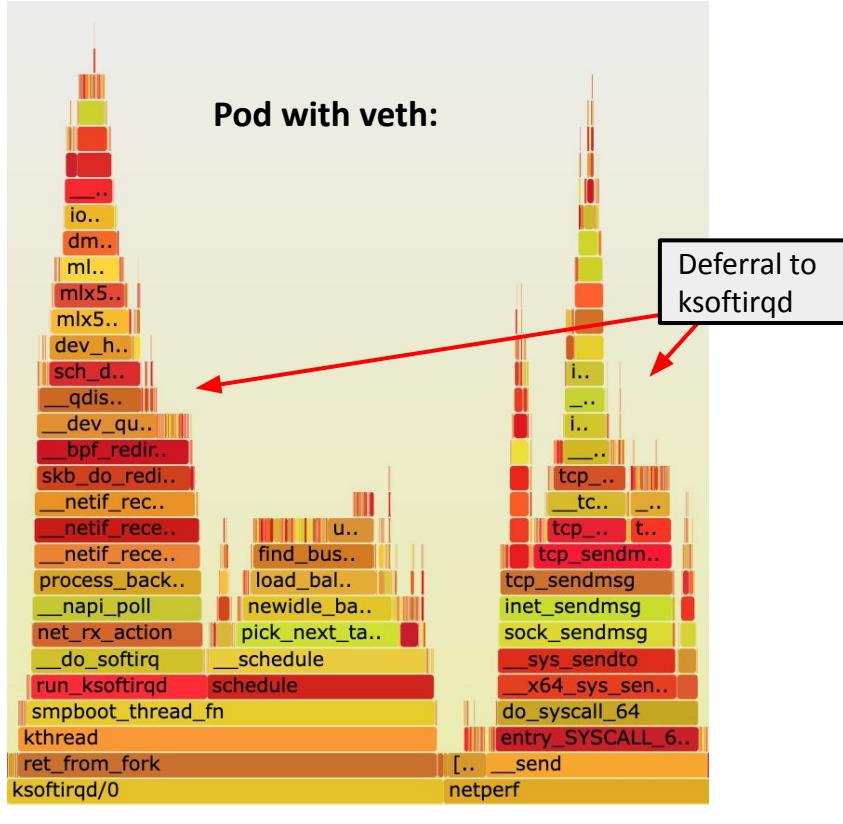


```
author      Daniel Borkmann <daniel@iogearbox.net> 2023-10-24 23:48:58 +0200
committer   Martin KaFai Lau <martin.lau@kernel.org> 2023-10-24 16:06:03 -0700
commit      35dfaad7188cdc043fde31709c796f5a692ba2bd (patch)
tree        53a88f1799ac38892434318a47278de4a255bc19
parent      42d31dd601fa43b9afdf069d1ba410b2306a4c76 (diff)
download    linux-35dfaad7188cdc043fde31709c796f5a692ba2bd.tar.gz
```

### netkit, bpf: Add bpf programmable net device

This work adds a new, minimal BPF-programmable device called "netkit" (former PoC code-name "meta") we recently presented at LSF/MM/BPF. The core idea is that BPF programs are executed within the drivers xmit routine and therefore e.g. in case of containers/Pods moving BPF processing closer to the source.

# veth vs netkit: Backlog Queue



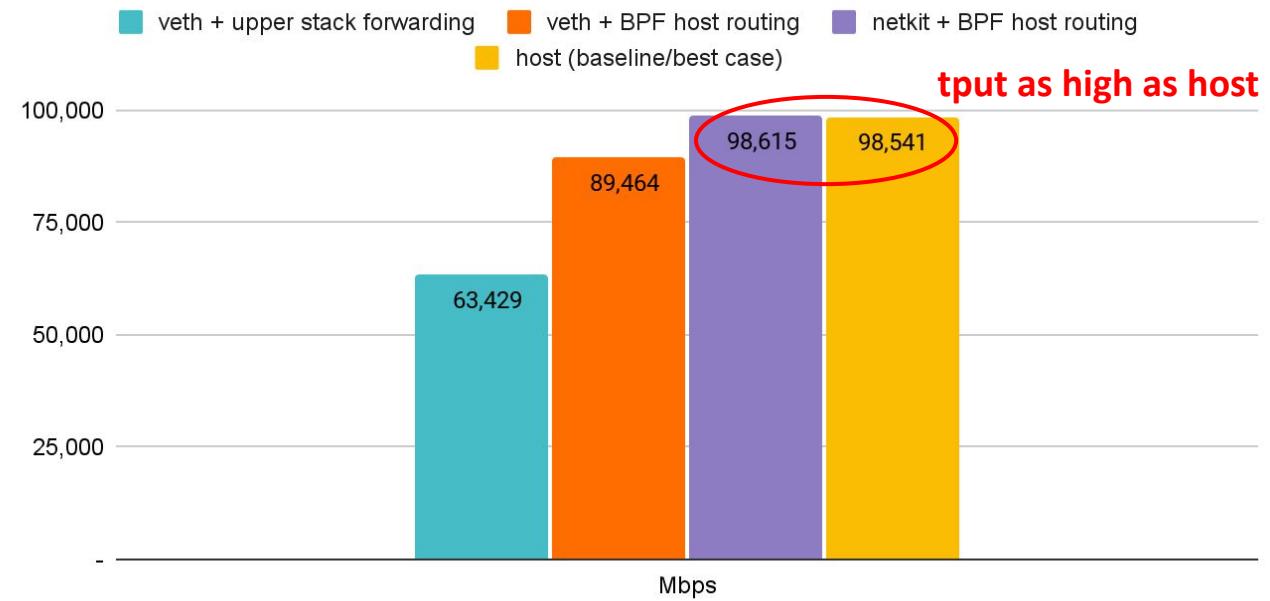


# Cilium's Pod Datapath: Making Pods as fast as Host

## Building Blocks:

- BPF kube-proxy replacement
- BPF Host Routing
- BIG TCP for GRO/GSO/TSO
- tcx-based BPF datapath layer
- netkit devices for Pods

TCP stream single flow Pod to Pod over wire, 8k MTU (higher is better)



\* 8264 MTU for data page alignment in GRO

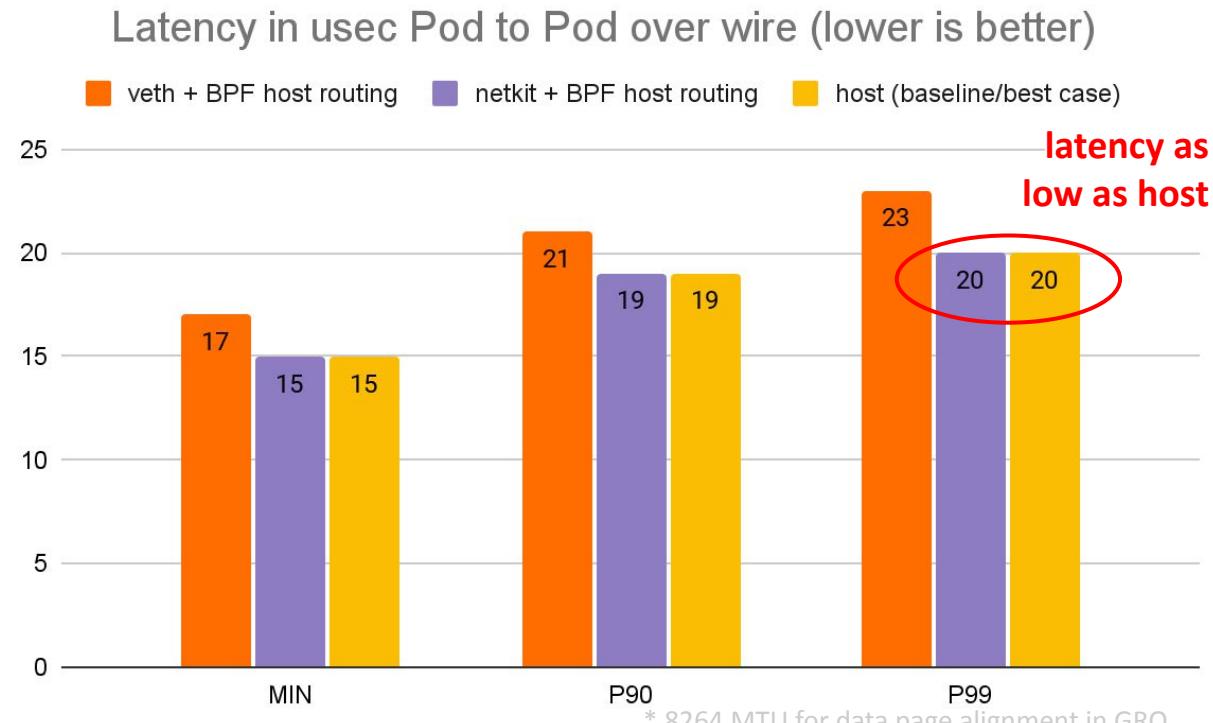
Back to back: AMD Ryzen 9 3950X @ 3.5 GHz, 128G RAM @ 3.2 GHz, PCIe 4.0, ConnectX-6 Dx, mlx5 driver, striding mode, LRO off, 8264 MTU  
Receiver: taskset -a -c <core> tcp\_mmap -s (non-zero-copy mode), Sender: taskset -a -c <core> tcp\_mmap -H <dst host>



# Cilium's Pod Datapath: Making Pods as fast as Host

## Building Blocks:

- BPF kube-proxy replacement
- BPF Host Routing
- BIG TCP for GRO/GSO/TSO
- tcx-based BPF datapath layer
- netkit devices for Pods



Back to back: AMD Ryzen 9 3950X @ 3.5 GHz, 128G RAM @ 3.2 GHz, PCIe 4.0, ConnectX-6 Dx, mlx5 driver, striding mode, LRO off  
netperf -t TCP\_RR -H <remote pod> -- -O MIN\_LATENCY,P90\_LATENCY,P99\_LATENCY,THROUGHPUT

CONTAINERS / EBPF / LINUX / NETWORKING

## Netkit to Network a Million Containers for ByteDance

Built on eBPF, netkit offers a swifter alternative to Virtual Ethernet for container networking, ByteDance engineers have concluded.

Jan 29th, 2025 6:00am by [Joab Jackson](#)

CONTAINERS / EBPF / LINUX / NETWORKING

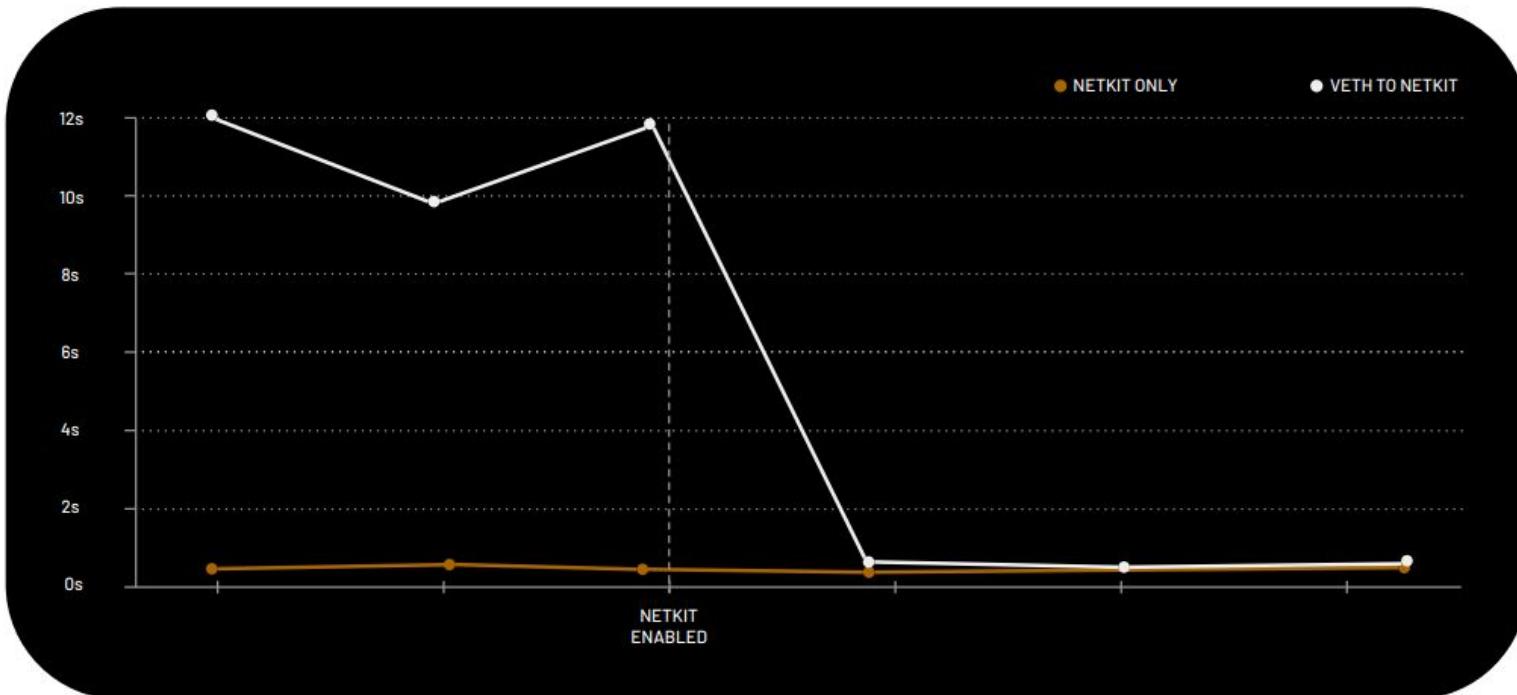
## Netkit to Network a Million Containers for ByteDance

Built on eBPF, netkit offers a swifter alternative to Virtual Ethernet for container networking, ByteDance engineers have concluded.

Jan 29th, 2025 6:00am by [Joab Jackson](#)

“Good news, we have finished the Proof of Concept (PoC) and netkit provided a 12% increase in CPS (Connections Per Seconds) compared with veth (incredible!). And now we plan to use netkit in our DC as much as possible.” - Chen Tang, ByteDance

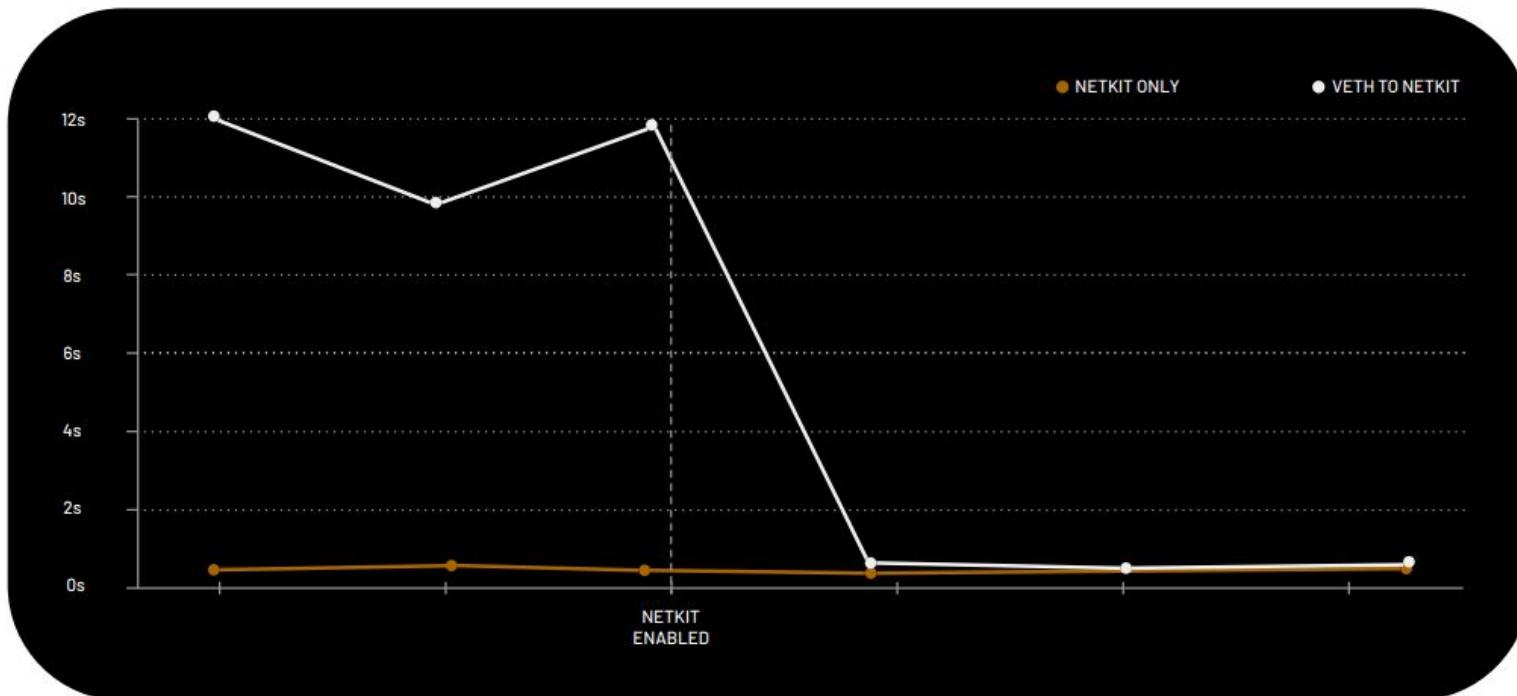
# netkit adoption outside of Cilium



100GiB single stream TCP @ 8K MTU

Meta (fleet-wide rollout): P99 reduction from 12s to 0.1s

# netkit adoption outside of Cilium



100GiB single stream TCP @ 8K MTU

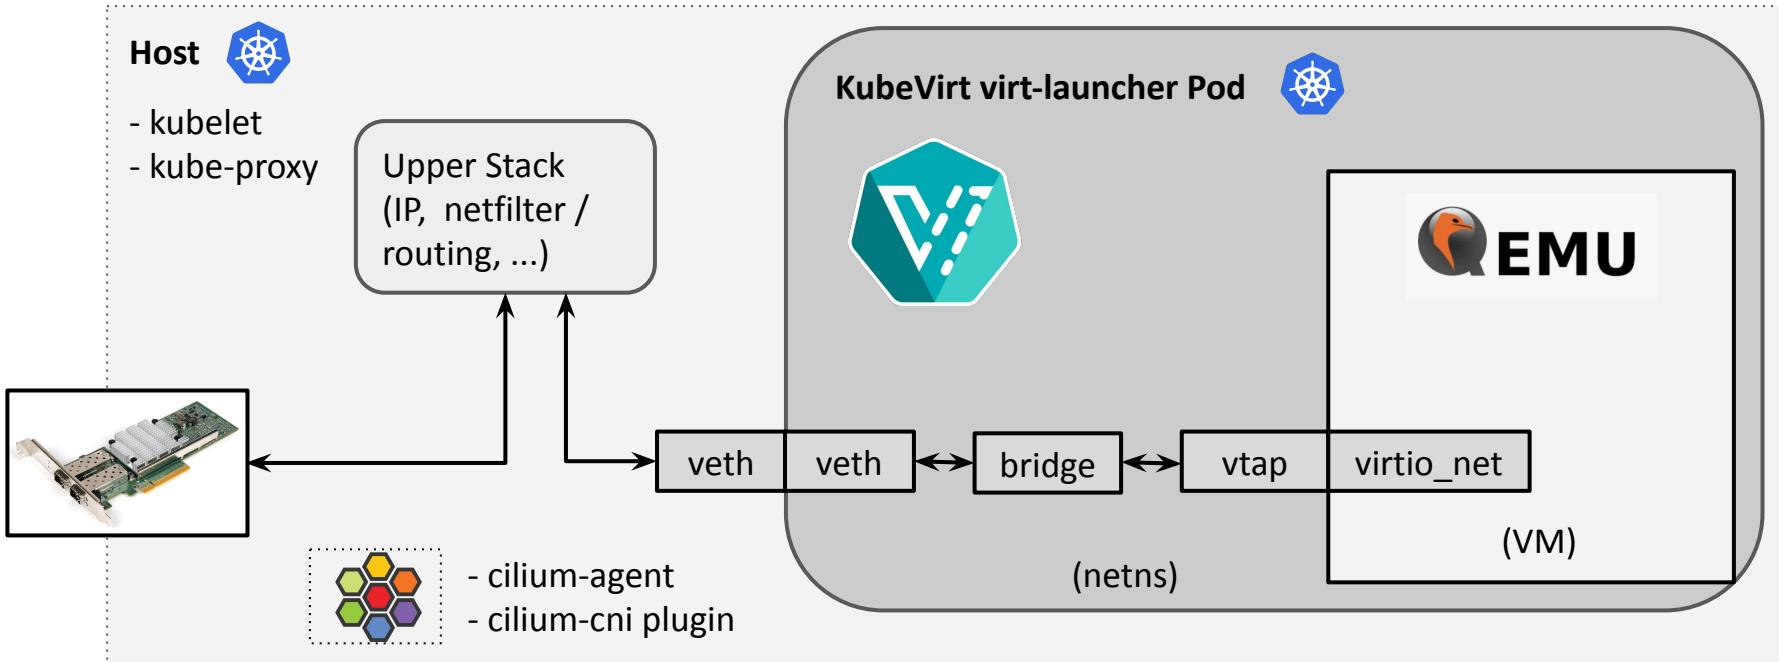
Meta (fleet-wide rollout): P99 reduction from 12s to 0.1s

"With netkit, network namespaces can be used without incurring a performance penalty." - Mike Willard, Meta

Could netkit also help in terms of VM networking?

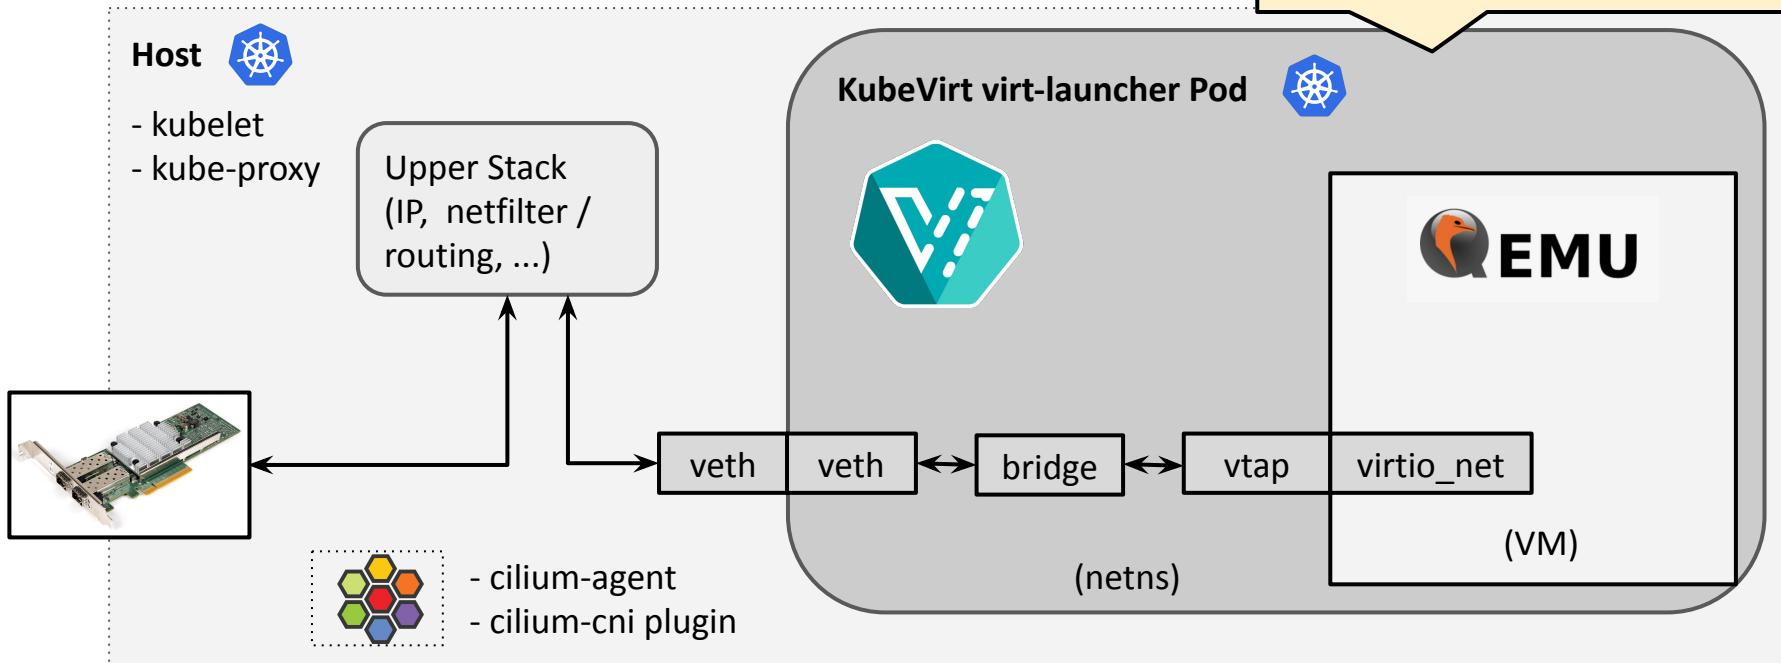


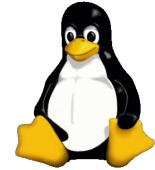
# Standard KubeVirt Pod Datapath:



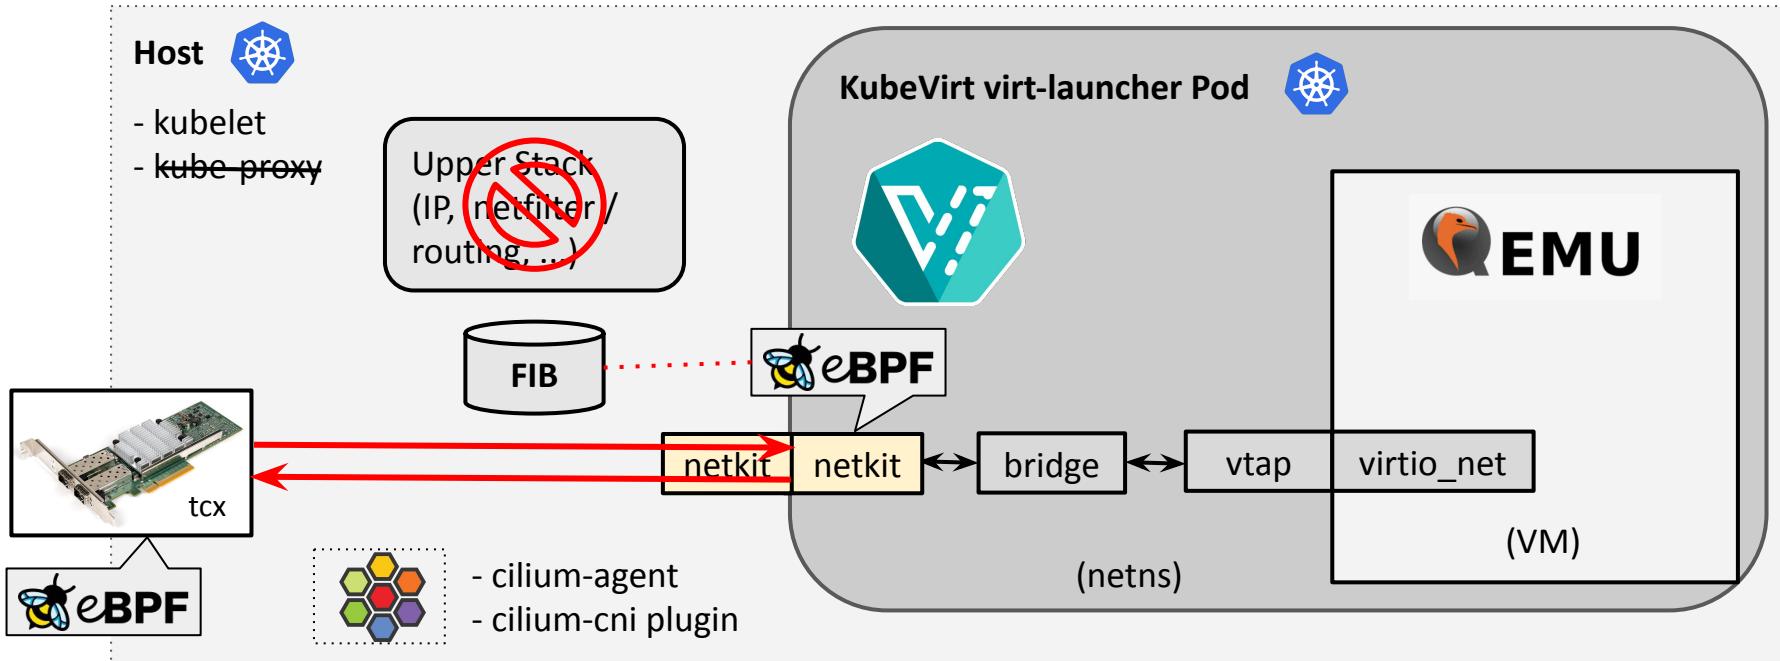
# Standard KubeVirt Pod Datapath:

Fyi: virt-launcher Pod provides ...  
- cgroups  
- namespaces  
... for hosting the VM process





# Cilium's KubeVirt Pod Datapath:





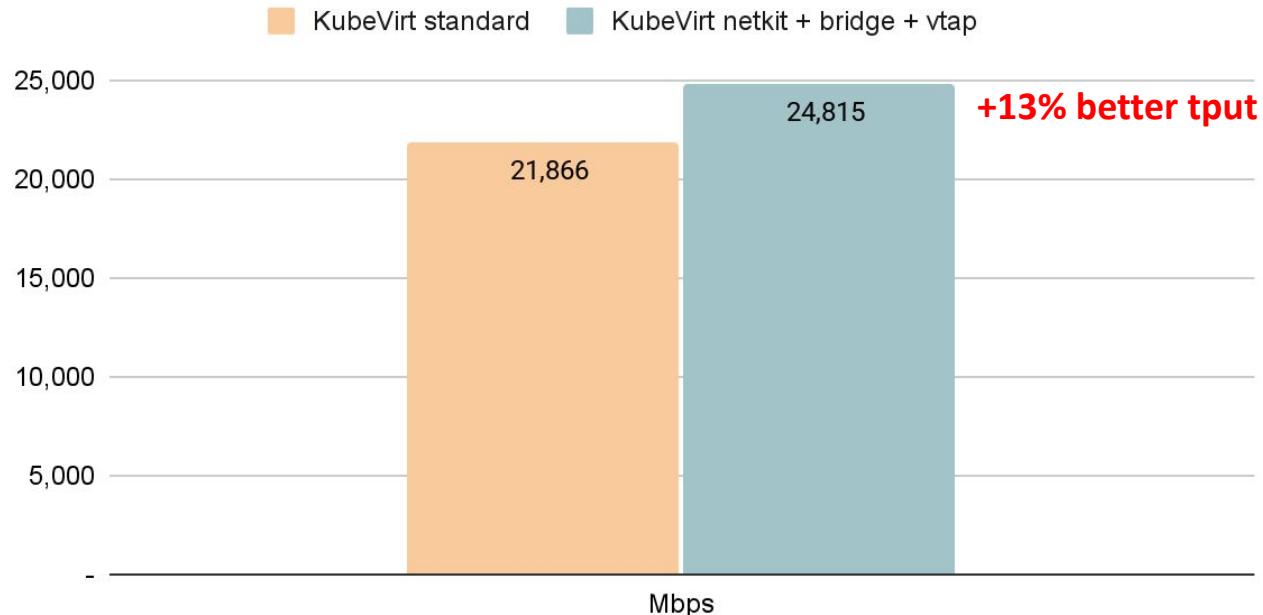
# Cilium's KubeVirt Pod Datapath:

## Building Blocks:

- BPF kube-proxy replacement
- BPF Host Routing
- BIG TCP for GRO/GSO/TSO
- tcx-based BPF datapath layer
- netkit devices for Pods

**(... but still KubeVirt with bridge + vtap)**

TCP stream single flow host to VM over wire  
1500 MTU (higher is better)



Back to back: AMD Ryzen 9 3950X @ 3.5 GHz, 128G RAM @ 3.2 GHz, PCIe 4.0, ConnectX-6 Dx, mlx5 driver, striding mode, LRO off, 1500 MTU  
netperf -t TCP\_RR -H <VM IP> --O MIN\_LATENCY,P90\_LATENCY,P99\_LATENCY,THROUGHPUT



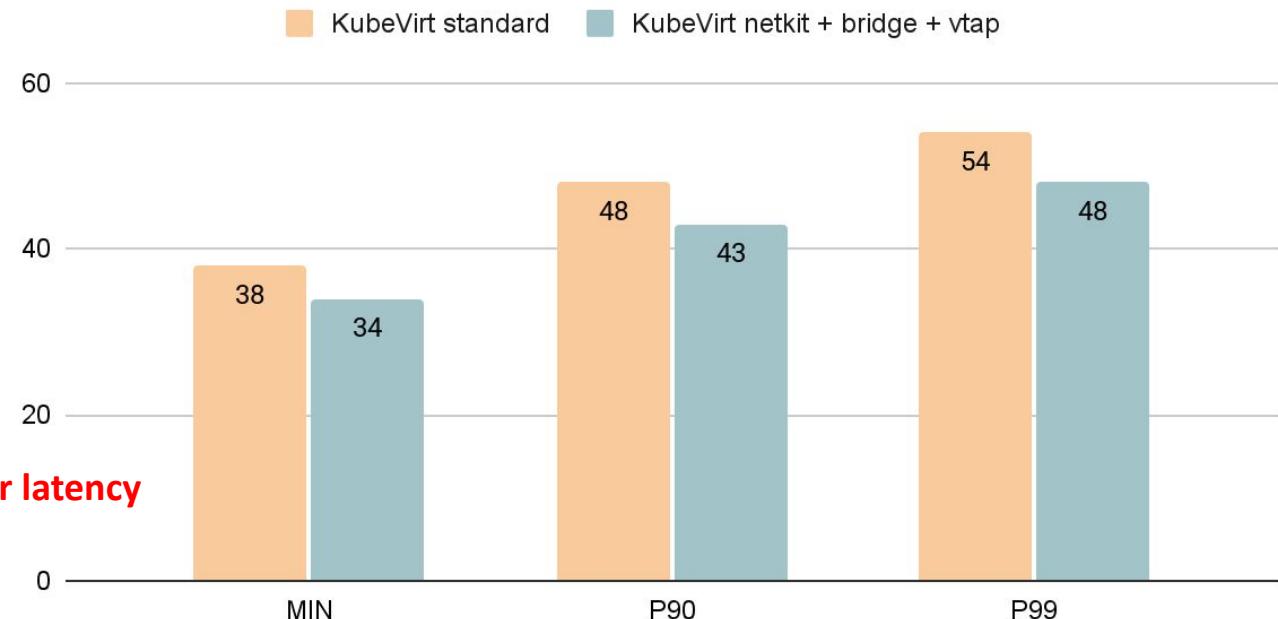
# Cilium's KubeVirt Pod Datapath:

## Building Blocks:

- BPF kube-proxy replacement
- BPF Host Routing
- BIG TCP for GRO/GSO/TSO
- tcx-based BPF datapath layer
- netkit devices for Pods

(... but still KubeVirt with bridge + vtap)

Latency in usec host to VM over wire  
(lower is better)

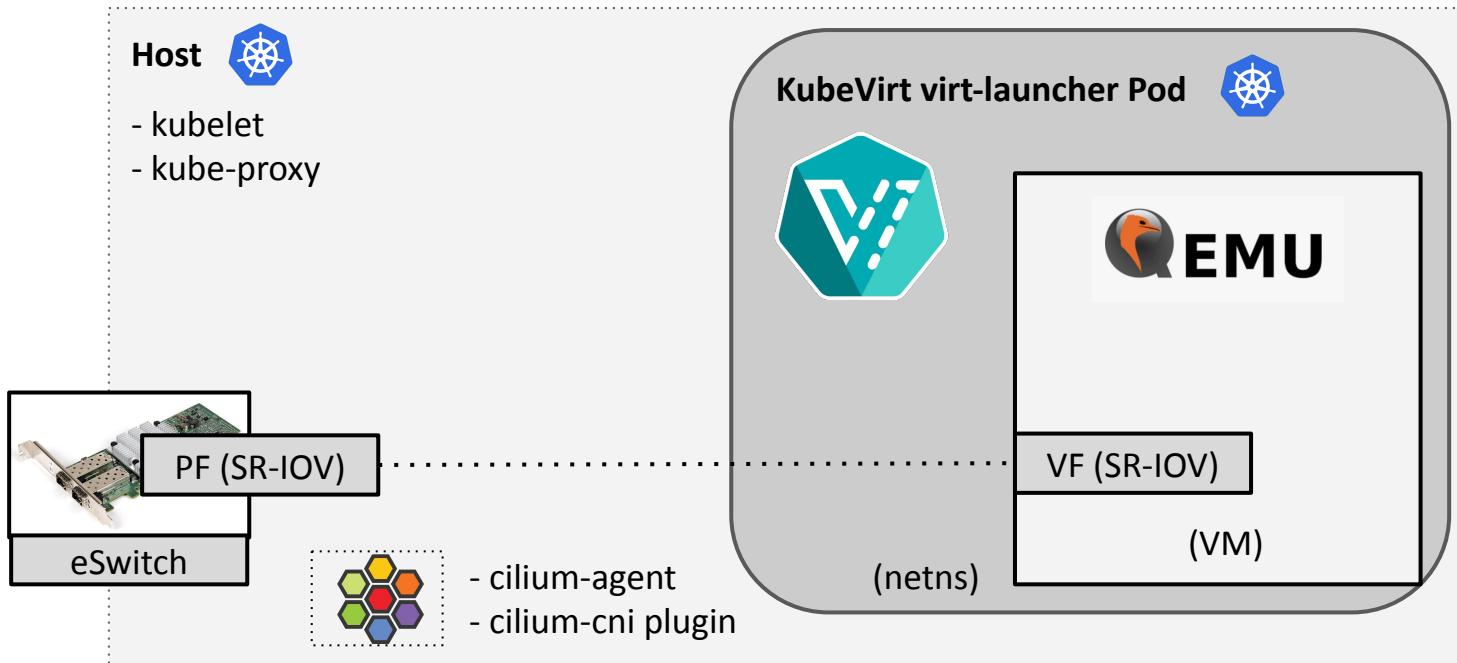


Back to back: AMD Ryzen 9 3950X @ 3.5 GHz, 128G RAM @ 3.2 GHz, PCIe 4.0, ConnectX-6 Dx, mlx5 driver, striding mode, LRO off, 1500 MTU  
netperf -t TCP\_RR -H <VM IP> --O MIN\_LATENCY,P90\_LATENCY,P99\_LATENCY,THROUGHPUT

Removal of the netns overhead helps a bit,  
but the main overhead is still on the VM side.

What about fundamentally challenging the  
bridge & vtap status quo?

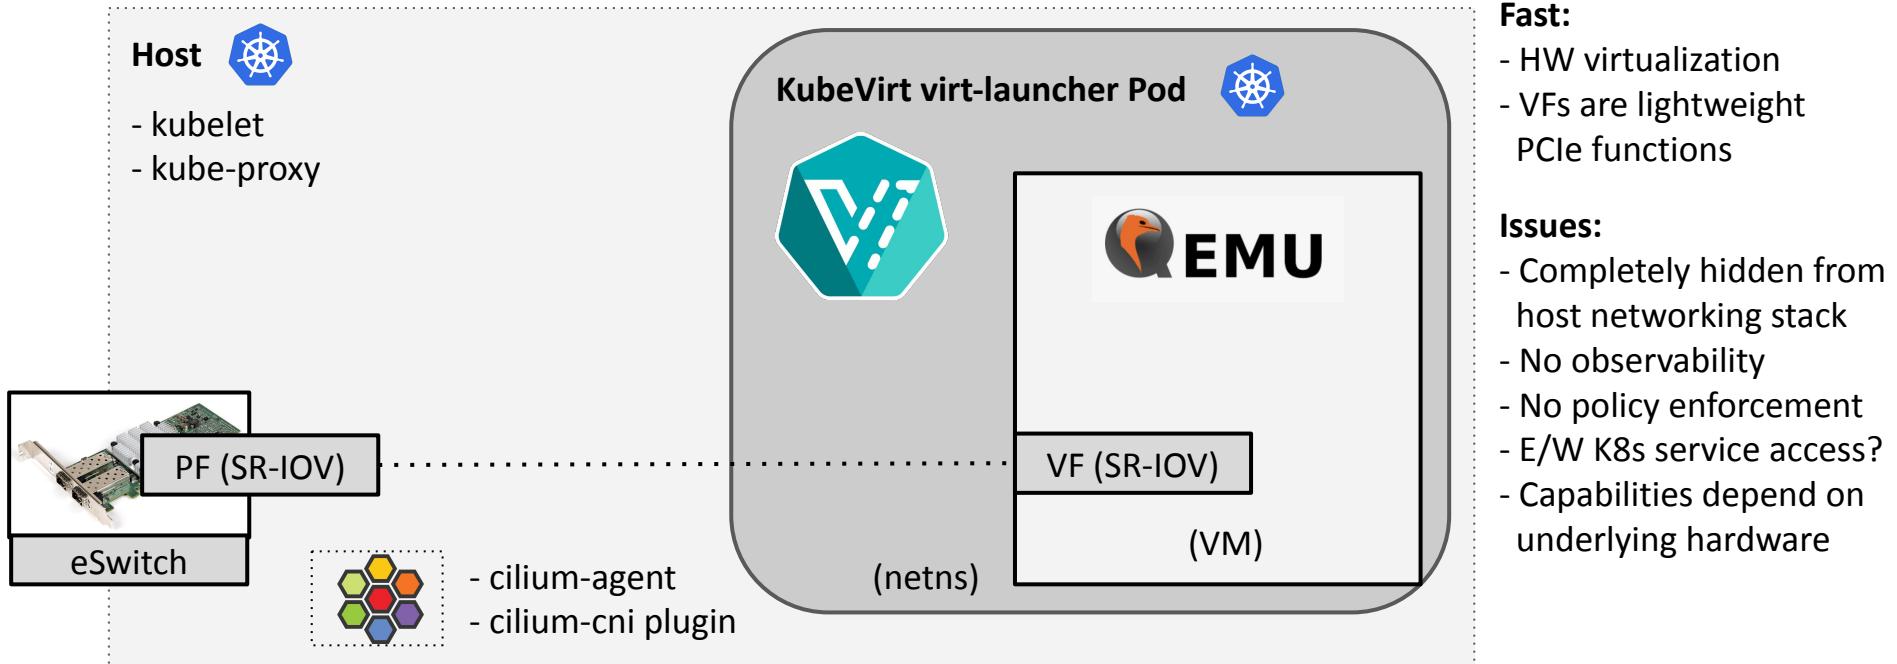
# What about SR-IOV?



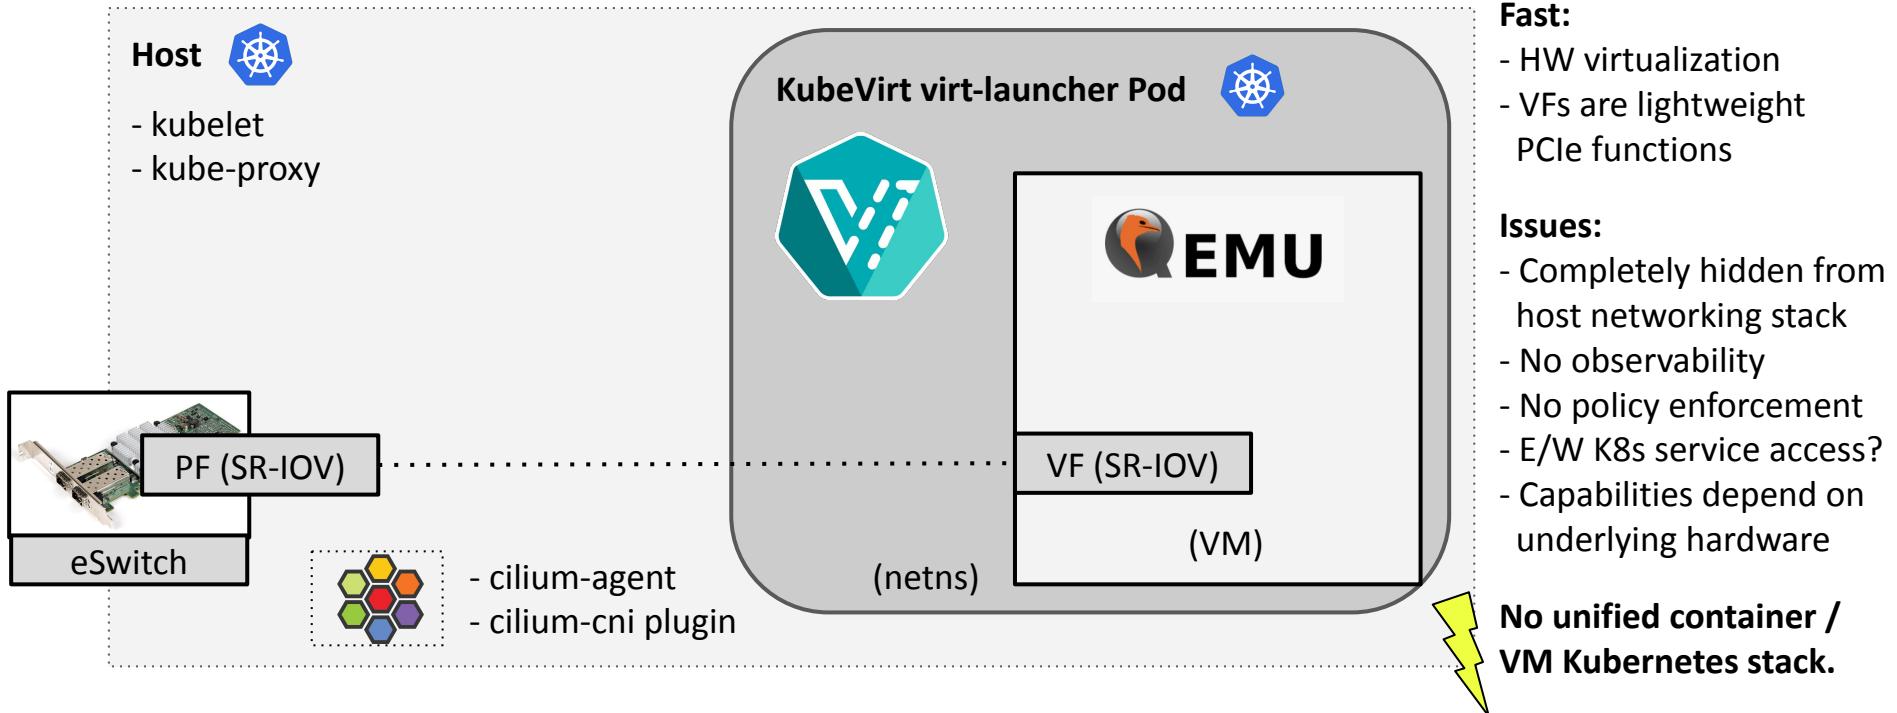
## Fast:

- HW virtualization
- VFs are lightweight PCIe functions

# What about SR-IOV?



# What about SR-IOV?



# Other alternatives?

 QEMU /  QEMU / Commits / cb039ef3

---

Commit cb039ef3  authored 1 year ago by  Ilya Maximets Committed by Jason Wang 1 year ago

---

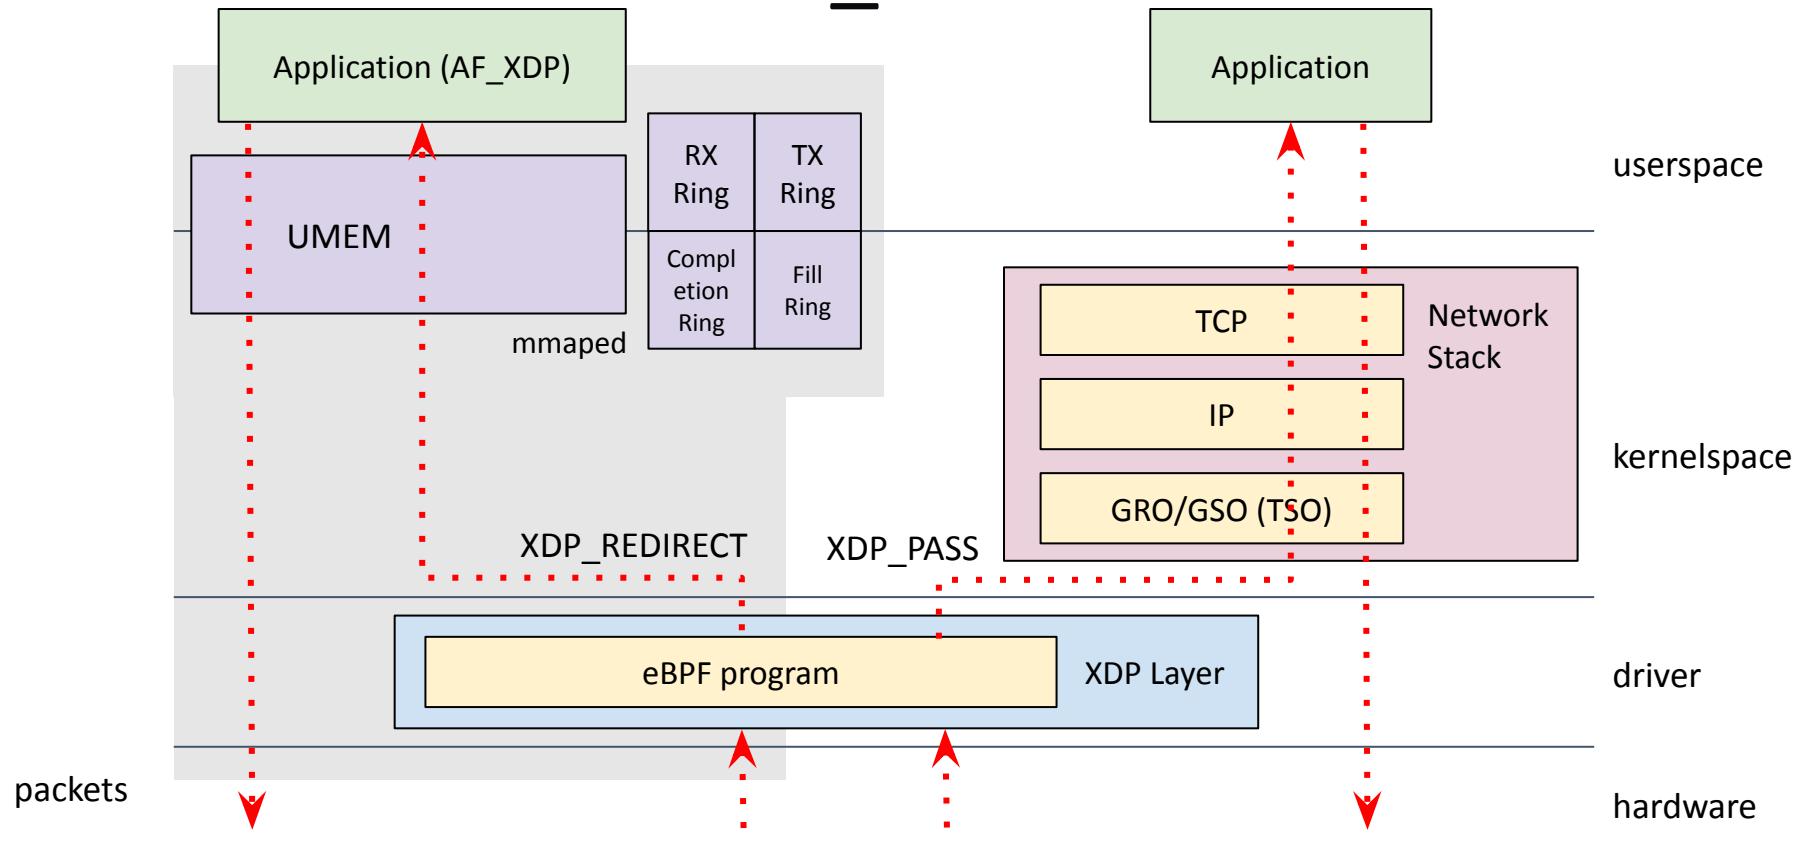
## net: add initial support for AF\_XDP network backend

AF\_XDP is a network socket family that allows communication directly with the network device driver in the kernel, bypassing most or all of the kernel networking stack. In the essence, the technology is pretty similar to netmap. But, unlike netmap, AF\_XDP is Linux-native and works with any network interfaces without driver modifications. Unlike vhost-based backends (kernel, user, vdpa), AF\_XDP doesn't require access to character devices or unix sockets. Only access to the network interface itself is necessary.

# AF\_XDP

- Zero-copy mechanism for high-performance networking
- Takes advantage of XDP datapath
- Native kernel integration, no out-of-tree drivers/vendor lock-in
- Original PoC (queue bifurcation proposal) which then led to the development of AF\_XDP was actually targeted at QEMU

# AF\_XDP



There's a catch: KubeVirt launches QEMU in its own Pod (netns). No AF\_XDP support in that netns given it needs to attach to a phys device.



# Deep Dive: netkit & Hardware Queue Binding for Pods

**netkit extension for *native* zero-copy performance inside Pods:**

- Physical NICs have multiple RX/TX queues, they can have multiple RSS contexts where traffic can be steered to

# Deep Dive: netkit & Hardware Queue Binding for Pods



**netkit extension for *native* zero-copy performance inside Pods:**

- Physical NICs have multiple RX/TX queues, they can have multiple RSS contexts where traffic can be steered to
- The default RSS context can be reduced and a new one sliced to steer traffic based on rules (e.g. MAC)

# Deep Dive: netkit & Hardware Queue Binding for Pods



**netkit extension for *native* zero-copy performance inside Pods:**

- Physical NICs have multiple RX/TX queues, they can have multiple RSS contexts where traffic can be steered to
- The default RSS context can be reduced and a new one sliced to steer traffic based on rules (e.g. MAC)
- netkit can then be created with numrxqueues > 1 and bind a physical to a virtual queue

# Deep Dive: netkit & Hardware Queue Binding for Pods



**netkit extension for *native* zero-copy performance inside Pods:**

- Physical NICs have multiple RX/TX queues, they can have multiple RSS contexts where traffic can be steered to
- The default RSS context can be reduced and a new one sliced to steer traffic based on rules (e.g. MAC)
- netkit can then be created with numrxqueues > 1 and bind a physical to a virtual queue
- The netkit device is moved into the target Pod (netns) and applications can bind against the netkit device + queue



# Deep Dive: netkit & Hardware Queue Binding for Pods

**netkit extension for *native* zero-copy performance inside Pods:**

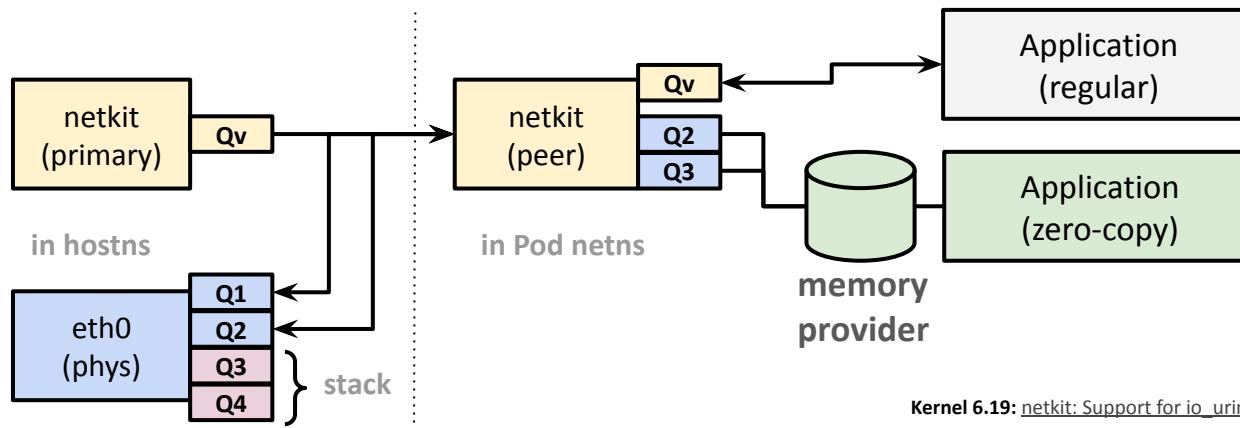
- Physical NICs have multiple RX/TX queues, they can have multiple RSS contexts where traffic can be steered to
- The default RSS context can be reduced and a new one sliced to steer traffic based on rules (e.g. MAC)
- netkit can then be created with numrxqueues > 1 and bind a physical to a virtual queue
- The netkit device is moved into the target Pod (netns) and applications can bind against the netkit device + queue
- This opens up native speed for AF\_XDP zero-copy, io\_uring & devmem zero-copy



# Deep Dive: netkit & Hardware Queue Binding for Pods

netkit extension for *native* zero-copy performance inside Pods:

- Physical NICs have multiple RX/TX queues, they can have multiple RSS contexts where traffic can be steered to
- The default RSS context can be reduced and a new one sliced to steer traffic based on rules (e.g. MAC)
- netkit can then be created with numrxqueues > 1 and bind a physical to a virtual queue
- The netkit device is moved into the target Pod (netns) and applications can bind against the netkit device + queue
- This opens up native speed for AF\_XDP zero-copy, io\_uring & devmem zero-copy

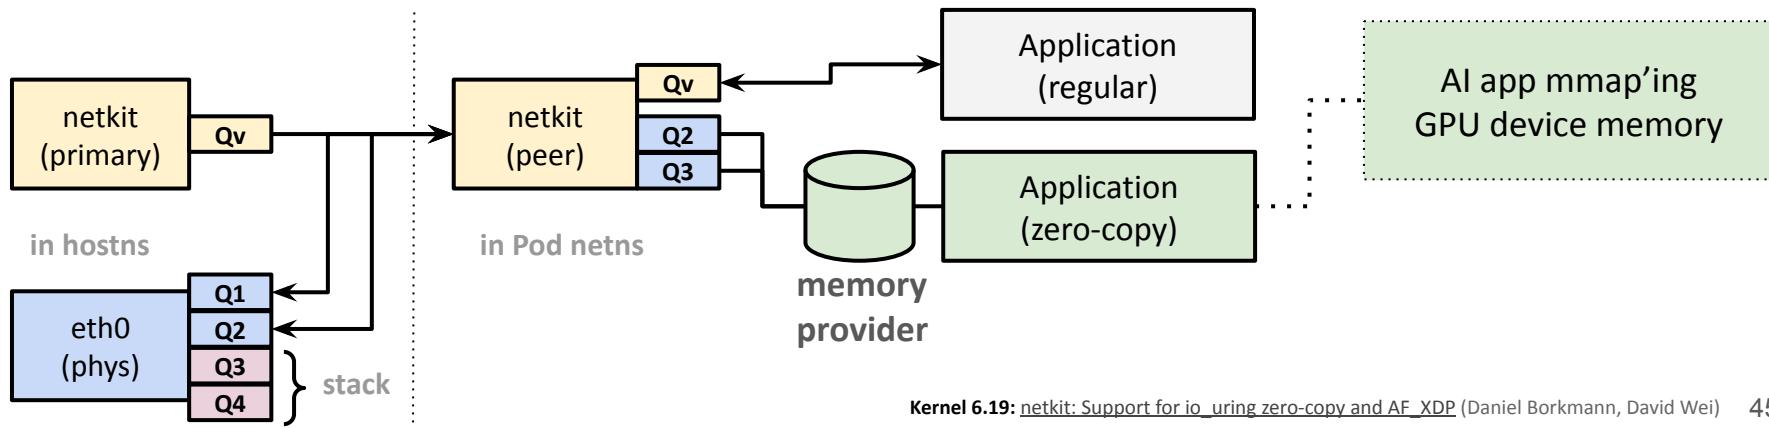


# Deep Dive: netkit & Hardware Queue Binding for Pods



netkit extension for *native* zero-copy performance inside Pods:

- Physical NICs have multiple RX/TX queues, they can have multiple RSS contexts where traffic can be steered to
- The default RSS context can be reduced and a new one sliced to steer traffic based on rules (e.g. MAC)
- netkit can then be created with numrxqueues > 1 and bind a physical to a virtual queue
- The netkit device is moved into the target Pod (netns) and applications can bind against the netkit device + queue
- This opens up native speed for AF\_XDP zero-copy, io\_uring & devmem zero-copy

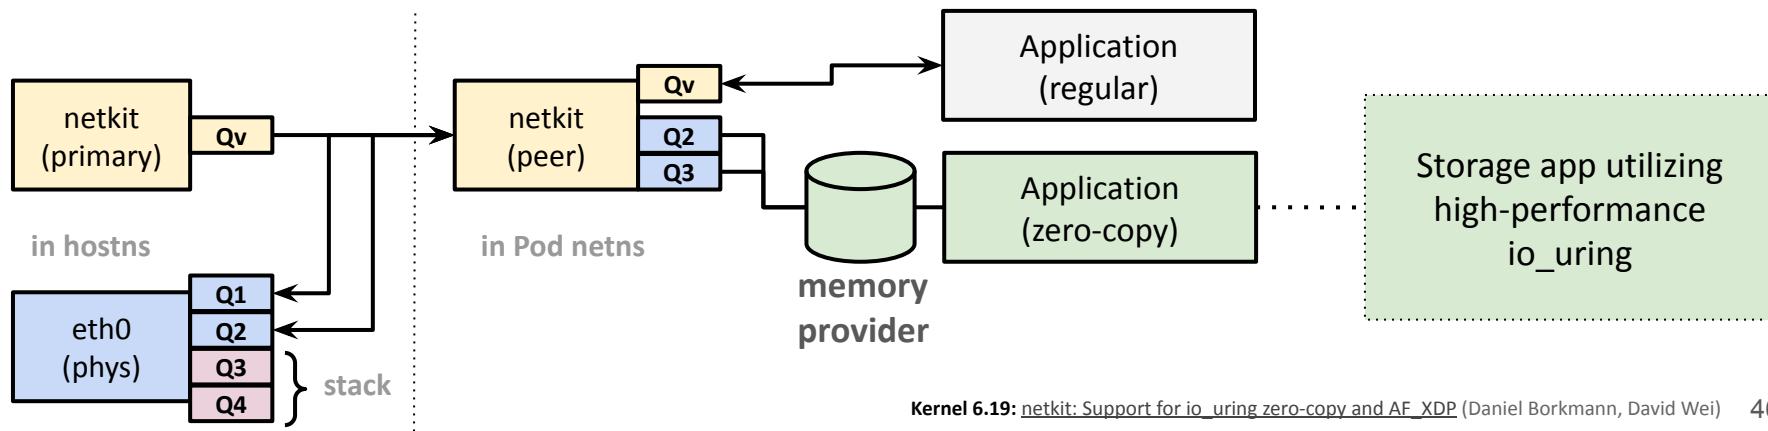




# Deep Dive: netkit & Hardware Queue Binding for Pods

netkit extension for *native* zero-copy performance inside Pods:

- Physical NICs have multiple RX/TX queues, they can have multiple RSS contexts where traffic can be steered to
- The default RSS context can be reduced and a new one sliced to steer traffic based on rules (e.g. MAC)
- netkit can then be created with numrxqueues > 1 and bind a physical to a virtual queue
- The netkit device is moved into the target Pod (netns) and applications can bind against the netkit device + queue
- This opens up native speed for AF\_XDP zero-copy, io\_uring & devmem zero-copy

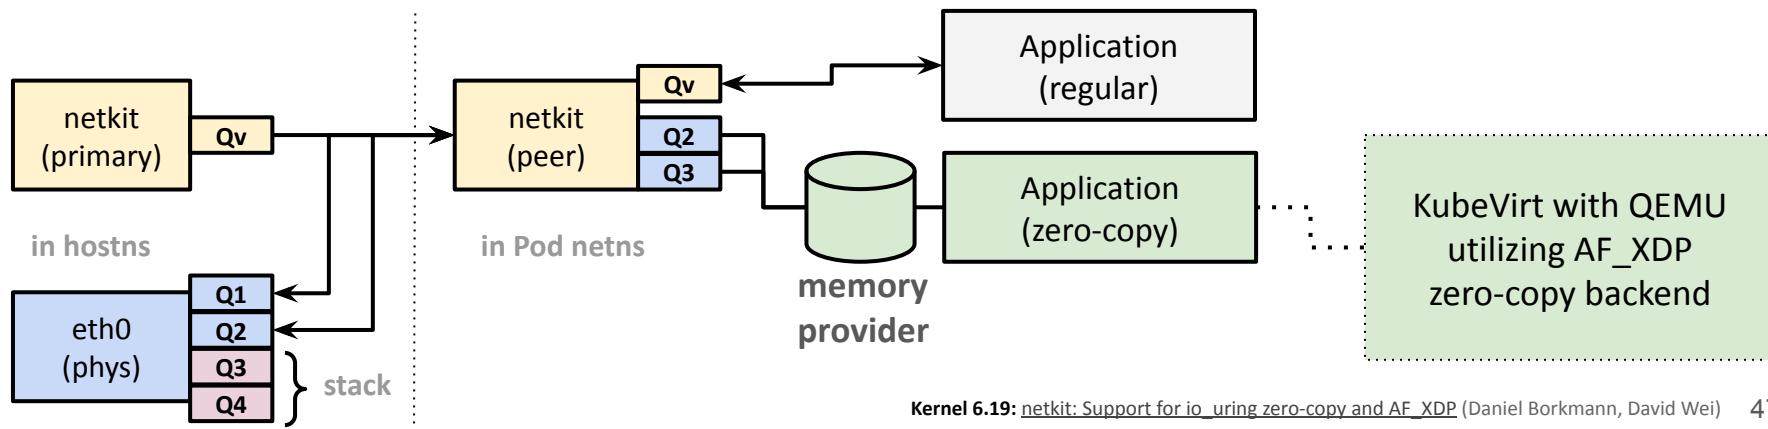




# Deep Dive: netkit & Hardware Queue Binding for Pods

netkit extension for *native* zero-copy performance inside Pods:

- Physical NICs have multiple RX/TX queues, they can have multiple RSS contexts where traffic can be steered to
- The default RSS context can be reduced and a new one sliced to steer traffic based on rules (e.g. MAC)
- netkit can then be created with numrxqueues > 1 and bind a physical to a virtual queue
- The netkit device is moved into the target Pod (netns) and applications can bind against the netkit device + queue
- This opens up native speed for AF\_XDP zero-copy, io\_uring & devmem zero-copy



# Same Queue Binding, but Different Zero-Copy ‘Beasts’...

## AF\_XDP zero-copy

- Full control of everything, e.g. useful for DPDK-like applications in user space
- Build-your-own load balancer, gateway, firewall, *QEMU* has a backend
- Upper kernel stack is bypassed, we reuse the driver, netdevice, XDP

## io\_uring with TCP zero-copy support

- The industry is now shifting towards header/data split for high-end NICs (bnxt/mlx5/gve)
- Data (payload) pages mmap’ed with user space application or GPU device memory
- Upper kernel stack is *reused*, not bypassed - packet traverses up the normal stack

# Results: 4096 MTU

[ [kernel-recipes](#) ]

MTU	memcmp	Engine	BW	Gain	Net CPU busy%	Net CPU softirq%
4096	✓	epoll	66.9 Gbps		24.8%	23.7%
4096	✓	io_uring ZC	92.2 Gbps	+37.8%	36.6%	35.0%
4096	✗	epoll	82.2 Gbps		33.2%	32.6%
4096	✗	io_uring ZC	116.2 Gbps	+41.4%	48.2%	46.6%

(beyond 100 Gbit/s for a single TCP stream requires zero-copy)

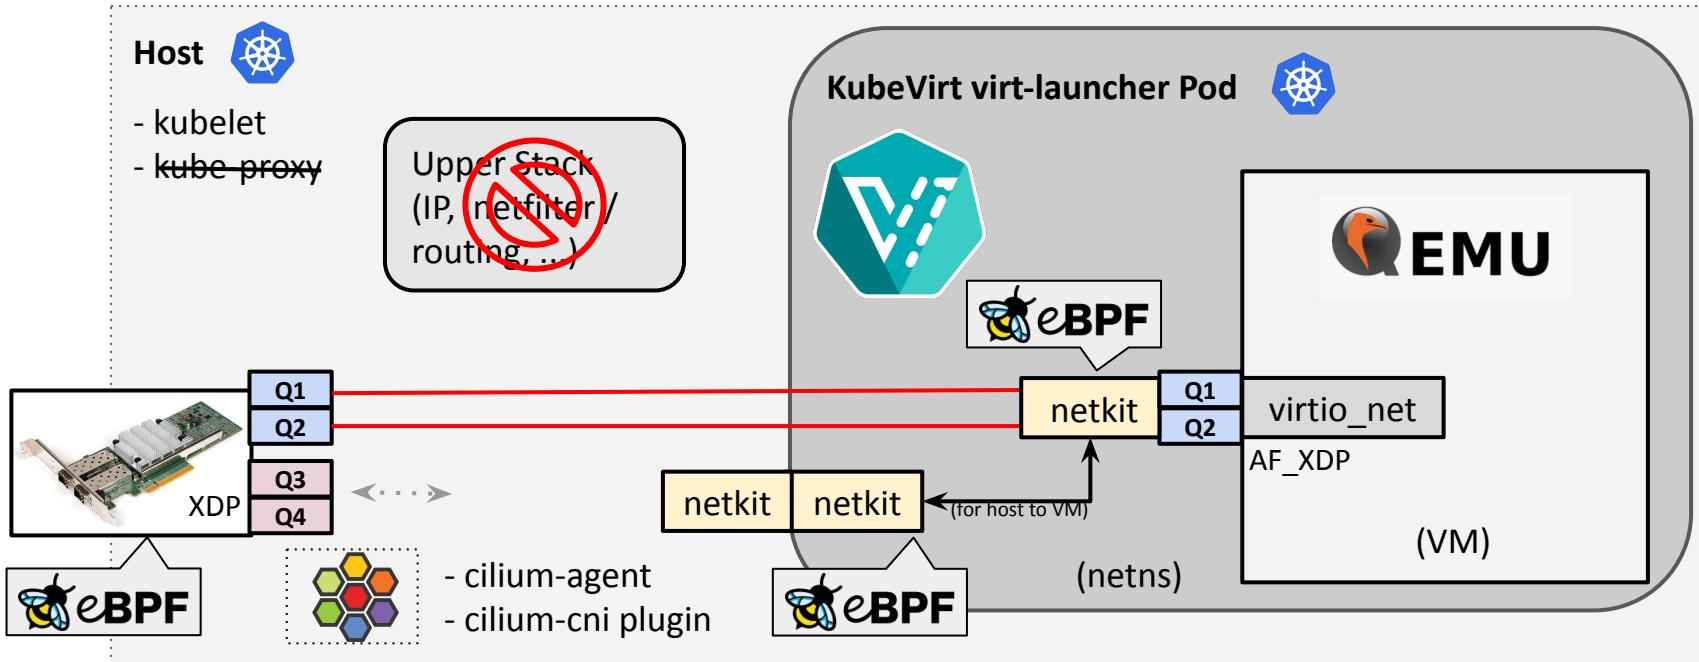
## io\_uring with TCP zero-copy support

- The industry is now shifting towards header/data split for high-end NICs (bnxt/mlx5/gve)
- Data (payload) pages mmap'ed with user space application or GPU device memory
- Upper kernel stack is *reused*, not bypassed - packet traverses up the normal stack

How can the netkit queue-binding & AF\_XDP  
be applied for KubeVirt Pods?



# Cilium's Future KubeVirt Pod Datapath:



Preliminary benchmarks with netkit + AF\_XDP  
for KubeVirt look promising!

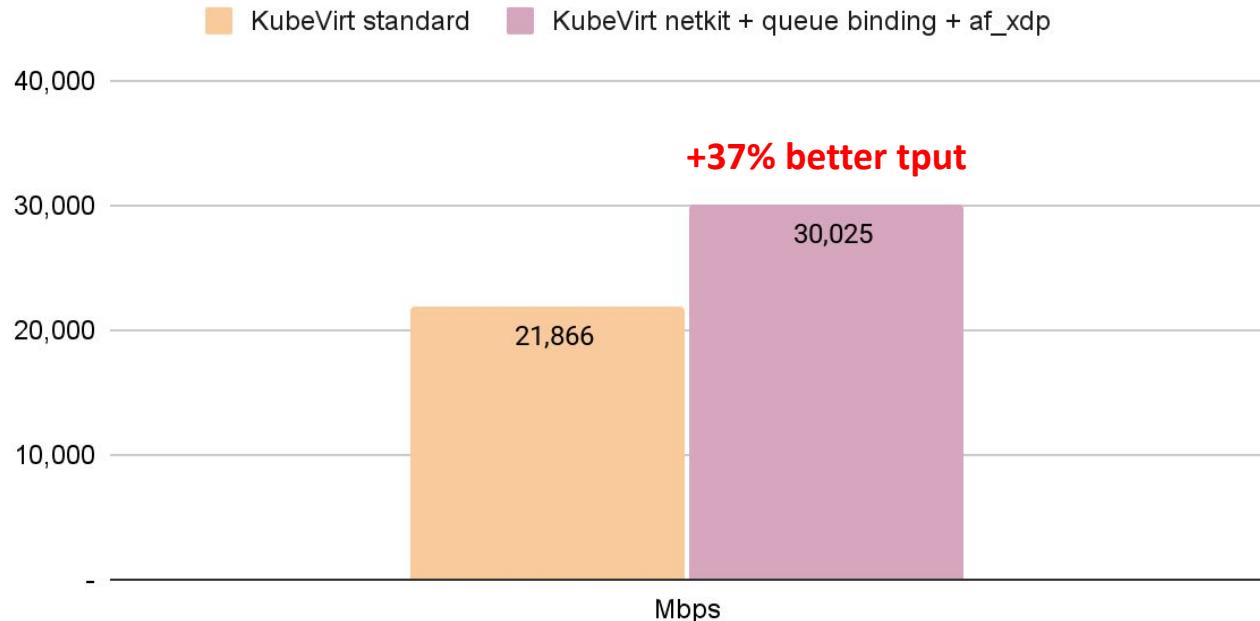


# Cilium's Future K8s KubeVirt Pod Datapath:

## Building Blocks:

- Cilium sets up new RSS context based on VM MAC
- Creates netkit device and binds netkit to phys queue
- Cilium KubeVirt plugin to launch QEMU with af-xdp bound to netkit

TCP stream single flow host to VM over wire  
1500 MTU (higher is better)



Back to back: AMD Ryzen 9 3950X @ 3.5 GHz, 128G RAM @ 3.2 GHz, PCIe 4.0, ConnectX-6 Dx, mlx5 driver, striding mode, LRO off, 1500 MTU  
netperf -t TCP\_RR -H <VM IP> --O MIN\_LATENCY,P90\_LATENCY,P99\_LATENCY,THROUGHPUT

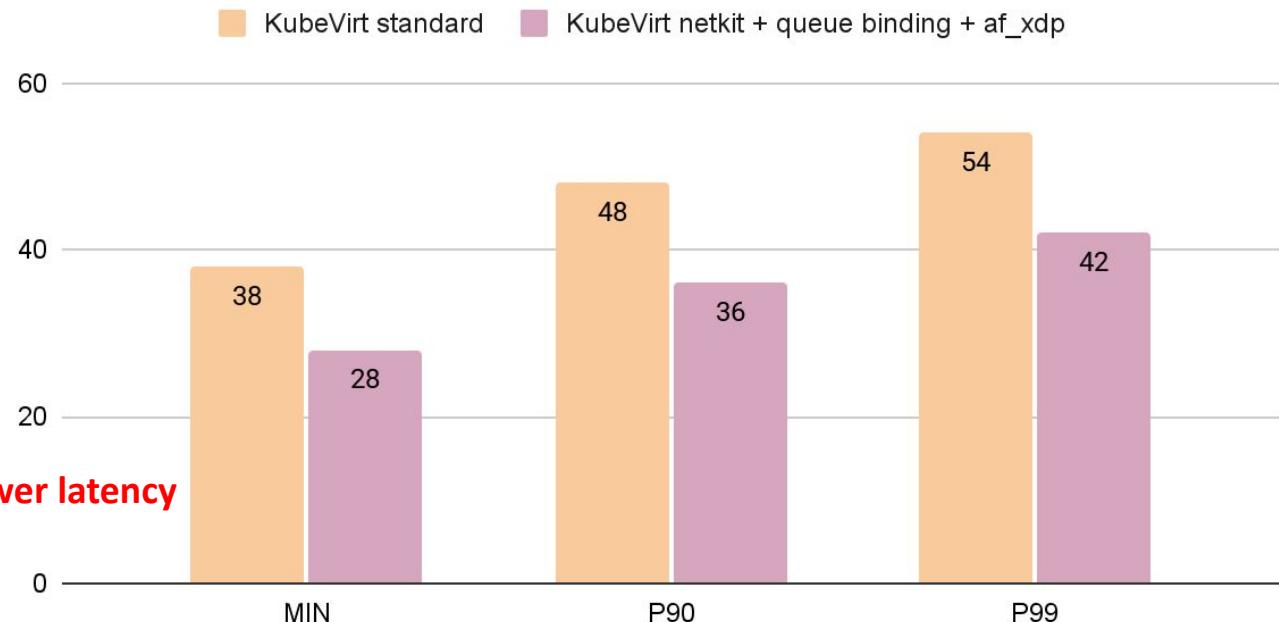


# Cilium's Future K8s KubeVirt Pod Datapath:

Latency in usec host to VM over wire  
(lower is better)

## Building Blocks:

- Cilium sets up new RSS context based on VM MAC
- Creates netkit device and binds netkit to phys queue
- Cilium KubeVirt plugin to launch QEMU with af-xdp bound to netkit



Back to back: AMD Ryzen 9 3950X @ 3.5 GHz, 128G RAM @ 3.2 GHz, PCIe 4.0, ConnectX-6 Dx, mlx5 driver, striding mode, LRO off, 1500 MTU  
netperf -t TCP\_RR -H <VM IP> --O MIN\_LATENCY,P90\_LATENCY,P99\_LATENCY,THROUGHPUT

# Ongoing & future work for KubeVirt acceleration

- XDP API overhaul in the Linux kernel:
  - Support for native multi-attach via bpf\_mprog infrastructure
  - Support for attaching BPF only to specific queues

# Ongoing & future work for KubeVirt acceleration

- XDP API overhaul in the Linux kernel:
  - Support for native multi-attach via bpf\_mprog infrastructure
  - Support for attaching BPF only to specific queues
- AF\_XDP extensions in the Linux kernel:
  - New AF\_XDP TX hook for BPF attachments
  - BPF xsk-map-less redirect to the target socket
  - AF\_XDP TSO offload support

# Ongoing & future work for KubeVirt acceleration

- XDP API overhaul in the Linux kernel:
  - Support for native multi-attach via bpf\_mprog infrastructure
  - Support for attaching BPF only to specific queues
- AF\_XDP extensions in the Linux kernel:
  - New AF\_XDP TX hook for BPF attachments
  - BPF xsk-map-less redirect to the target socket
  - AF\_XDP TSO offload support
- Further optimize QEMU's AF\_XDP integration:
  - Various low-hanging fruit: utilizing scatter/gather IO in backend
  - Bigger items further helping performance: full zero-copy into VM

Small demo: Cilium and KubeVirt with netkit's  
queue binding and AF\_XDP QEMU backend ...

```
+ yq -C .spec.template < vm.yaml | head -20
{
  "metadata": {
    "labels": {
      "kubevirt.io/vm": "vm-net-binding-netkit"
    },
    "annotations": {
      "cilium.io/bind-queue": "enp2s0f1np1/15",
      "cilium.io/mac": "18:c0:4d:77:88:99"
    }
  },
  "spec": {
    "domain": {
      "devices": {
        "interfaces": [
          {
            "name": "netkit",
            "binding": {
              "name": "netkit"
            }
          }
        ]
      }
    }
  }
}
```



# Acknowledgements

Ilya Maximets (Red Hat; QEMU's initial AF\_XDP backend integration)

Nikolay Aleksandrov (Nvidia; netkit co-creator)

Jordan Rife, Alasdair McWilliam (Google & Isovalent; Cilium netkit contributions)

David Wei, Jakub Kicinski, Stanislav Fomichev (Meta; io\_uring & kernel networking)

Martin Lau (Meta; netkit production rollout)

Chen Tang (Bytedance; netkit production rollout)

Björn Töpel, Magnus Karlsson, Maciej Fijalkowski (Rivos & Intel, AF\_XDP)

Cilium, netdev, BPF & QEMU communities



# Thank you! Questions?

[github.com/cilium/cilium](https://github.com/cilium/cilium)

[cilium.io](https://cilium.io)

[ebpf.io](https://ebpf.io)

[kubevirt.io](https://kubevirt.io)

Cilium & KubeVirt Instruqt Intro Lab

Cilium & netkit Tuning Guide

netkit Linux kernel upstream driver